

# Correctness-Gated, LLM-Driven Kernel Optimization on the Raspberry Pi 5’s Integrated GPU

Adam Munawar Rahman  
New York University  
New York, NY, USA  
msr541@nyu.edu

**Abstract**—Most edge AI deployments on single-board computers leave the integrated GPU idle. We present *Seppa*, a harness for optimizing GPU kernels with a large language model in the loop: the model proposes changes, and a finite-state machine compiles, verifies against physics gates, benchmarks, and decides keep or revert, refusing to benchmark any kernel that failed verification. On the flood-simulation stencil of an offline flood-guidance application for the Raspberry Pi 5, the harness produced a kernel 1.59x faster than the original; the machine re-measured and re-judged the result over the Model Context Protocol in five consecutive runs; each run also refused an out-of-order benchmark request for a gate-failing kernel. With CPU inference running concurrently the advantage widens to 2.20x, and the deployment sustains 9.6 tokens/s of language-model decode beside 883 simulation steps per second, an operating point that beats every measured CPU-only configuration on both axes.

**Index Terms**—GPU kernel optimization, LLM agents, correctness verification, edge computing, Vulkan

## I. INTRODUCTION

Raspberry Pi class boards are a common platform for edge AI in emergency-response and field settings, where power, connectivity, and budget are constrained [12] and inference must happen on-device once the network is gone [11]; a Pi 5 runs a quantized billion-parameter model on its CPU cores [8]. AI capability is usually added with dedicated hardware, an NPU HAT or a USB accelerator [13], which raises cost and, in a supply-constrained chip market [14], adds a procurement dependency.

Less attention goes to silicon already on the board: every Pi 5 ships a VideoCore VII GPU that compute workloads rarely touch, and during CPU inference it idles. Two questions decide that idle GPU’s worth: can it be made fast enough at a useful workload to be a co-processor, and can it run beside a busy CPU on a shared LPDDR bus? Sections 5 and 6 answer the first question, Section VII the second.

Making kernels fast on a sparsely documented mobile GPU, where tuning folklore from CUDA-class hardware mostly fails (Section II), is exactly the kind of work the field has begun handing to language models. Existing agentic optimizers, AutoKernel [15] among them, steer the agent with prompts and trust it to verify and revert honestly; the record shows what that trust costs (Section III). *Seppa* removes the trust. The model proposes, and nothing else; a finite-state machine built

on Burr [5] and served over the Model Context Protocol [6] owns compilation, verification, benchmarking, and the verdict, and refuses out-of-order requests.

The demonstration workload is the flood simulation of an offline flood-guidance application [7], which must run beside language-model narration on one board; its stencil, a neighbor-update grid kernel, is the optimization target, and its CPU-plus-GPU split is what we measure (Section IX describes the application).

This paper reports:

1. the optimization of a real flood stencil for V3D under three physics gates (correctness checks passed before any benchmark), to 1.59x, failed hypotheses included (Section V);
2. an action-space lesson: the first search plateaued because every winning change lived in the host program, outside what the machine let the model edit (Section V-C);
3. a machine-checked reproduction over MCP, repeated five times under a defined pass criterion, including the server’s refusal to benchmark a gate-failing kernel (Section VI);
4. an interference measurement of the deployed configuration, with a gate-verified CPU-only counterfactual that the GPU split beats on both axes (Section VII).

## II. BACKGROUND: THE RASPBERRY PI 5 GPU

The Pi 5 pairs four Cortex-A76 CPU cores with a VideoCore VII GPU (V3D 7.1), the block that also drives the desktop. There is no CUDA and no vendor compute toolchain; compute reaches the GPU only through Vulkan shaders, compiled by Mesa’s v3dv driver [9]. In this paper, kernel means a GLSL compute shader.

The limits are tight: at most 256 invocations and 16 KB of shared memory per workgroup, a fixed 16-lane SIMD width, no usable fp16, no matrix hardware. The best kernel we measured on this device sustains about 13 GFLOP/s in fp32; the four CPU cores manage about 5.5 GFLOP/s on comparable work. The GPU is a second, slower engine that happens to be free.

The strange part is the driver: when a shader wants more registers than exist, v3dv does not fail but recompiles with scheduling disabled, then unrolling disabled, then fewer




Fig. 1. The Seppa loop. Shaded states are the agent’s; dashed guard edges route compile and gate failures to `log_variant`, so benchmark requires a green verify.

threads, handing back whatever survives, sometimes several times slower.

### III. RELATED WORK

KernelBench [3], the standard measure of LLM kernel generation, conditions speedup on correctness by definition; the systems evaluated on it still cheat. CUDA-L1 [4] reports that 82 of its 250 RL-generated kernels (32.8%) exploited stream-timing loopholes to fake speedups of up to 18x, and containment took a reward checker, a database of known hacks, and forced stream synchronization. In every published case the fix is the same: verification the model cannot touch. Seppa builds that fix into the control flow, as a state the loop must pass through (Section IV). Chip-Chat [10] carried conversational hardware design to tapeout with a human engineer and a testbench closing the loop; we close it with a state machine instead.

### IV. THE SEPPA HARNESS

Seppa is a port of AutoKernel [15], an open-source kernel-optimization loop, onto a Burr finite-state machine, served over MCP by a wrapper (Theodosia) running on the Pi itself.

Fig. 1 shows the graph, with two guard edges: `compile -> log_variant` on compile failure, and `verify -> log_variant` on gate failure. The benchmark state is reachable only through a green verify. The agent’s entire interface is one MCP tool, `step(action, inputs)`, and the server constrains `action` to the graph’s legal next moves.

The port preserves AutoKernel’s loop discipline: one focused change per experiment, a baseline measured first, correctness and timing taken from the same execution, keep only on a gain of at least 1%, every experiment recorded. It moves one thing: the authority to decide. AutoKernel trusts its agent to run the benchmark and honestly revert failures; here the verdict is a deterministic action the model cannot reach. A run ends after three consecutive non-improvements on gate-passing variants, or at the budget. An audit of the port is in the repository (`docs/autokernel_fidelity.md`).

For the division of labor, we follow Chip-Chat’s disclosure convention. The author built the application and the harness, chose the physics gates, ran the hand-driven sweep of Section V-B, and supervised sessions. The language model,

TABLE I  
FALSIFICATION SWEEP: 400 STEPS AT 256x256 IN THE VKFLOOD2 HARNESS

| Variant                                                      | Hypothesis tested                      | Result                                     |
|--------------------------------------------------------------|----------------------------------------|--------------------------------------------|
| Packed vec2 state (water, surface) halving flux-pass loads   | Memory-op count is the bound           | 1.93 GFLOP/s, no change: falsified         |
| Fused single-dispatch step (flux buffer eliminated)          | Traffic and barriers are the bound     | 2.04 GFLOP/s, +5%: mostly falsified        |
| 2 cells per invocation (strip-mined dispatch)                | Fixed per-invocation cost is the bound | 2.40 GFLOP/s, +23%: supported              |
| 4 cells per invocation                                       | More strip is better                   | 2.35 GFLOP/s: falsified, register pressure |
| Fused + 2 cells per invocation ( <code>fused2s.comp</code> ) | The wins stack                         | 3.08 GFLOP/s, +59%                         |

Claude Sonnet 5 (Anthropic), driven through the Claude Code client at high reasoning effort, wrote the content of `hypothesize` and `implement` (kernel source and host-side parameters) and nothing else. The machine did everything evidential; no proposed kernel was hand-edited (transcripts in `docs/paper/artifacts/claude_sessions/`).

For the flood target, `verify` runs 400 steps at 256x256 and demands three things: normalized mean-square error (NMSE) below  $1e-3$  against a double-precision CPU reference, total water equal to injected rainfall, and maximum depth pooling inside the terrain basin; measured NMSE in practice is  $1.3e-9$ .

### V. CASE STUDY: THE FLOOD STENCIL

#### A. The Workload

The application’s simulation is a stencil; each cell updates from its neighbors, in two passes per time step, a flux pass and a height pass. The original implementation ran about 1,045 steps/s at 256x256 (1.51 GFLOP/s), and both shaders compile cleanly, which means they never stress the register allocator. The bottleneck was therefore something other than register pressure.

#### B. Falsification Sweep

Rather than guess the bottleneck, each hypothesis got its own variant and the loop measured each. The variants pack the state into wider loads, fuse the two passes into one dispatch, or strip-mine, meaning each thread updates a short vertical run of cells to amortize its launch cost. Every variant passed all three gates. Speeds are 400 steps at 256x256 on the parameterized `vkflood2` harness. Its host loop is faster than the original’s for identical kernels (about 1,345 steps/s against 1,045), so every comparison in this paper is within `vkflood2`. Table I shows the sweep.

The outcome surprised us. Halving the flux pass’s memory traffic changed nothing, and fusion, which deletes an intermediate buffer and a barrier, bought only 5%. Fixed per-invocation cost dominated: each invocation does so little arithmetic that launch overhead swamps it, and giving each

thread two cells beat every memory-system optimization we tried; four cells gave the gain back to register pressure. Two hardware details shaped the final kernel. Strips run vertically because that keeps a subgroup’s 16 lanes on adjacent addresses, and the kernel consumes each flux value as it is produced; holding all four alive fails register allocation. The shipped kernel is 1.59x at 256x256 (0.187 s against 0.297 s), 1.41x at 512x512, 1.18x at 1024x1024, and holds all gates over 4,000 steps.

### C. Action-Space Limits

An earlier session pointed the FSM at this stencil and came back empty; we nearly concluded the kernel was at its limit. The problem was the action space: that version of `implement` accepted shader text only, and every winning change in Table I lives outside the shader (packing changes the buffer layout, fusion deletes a host-loop pipeline, strip-mining changes the dispatch shape). Once `implement` took `{shader, height_shader, strip}`, the machine found and kept the fused strip-2 kernel from its own baseline (Section VI). A plateau reflects the search space before it reflects the hardware. The harness has also run against other kernels here: a GEMM (dense matrix multiply) went from 7.02 to 13.42 GFLOP/s over two rounds, and de-unrolling llama.cpp’s matrix-vector kernel gained 28% end-to-end decode.

## VI. MACHINE-CHECKED REPRODUCTION OVER MCP

Rather than ask a reader to trust Section V’s numbers, `drive_flood2_mcp.py` runs on a second computer and drives the Pi-resident FSM through the whole cycle via the `step` tool.

The FSM measures its own baseline (1,346.8 steps/s on 2026-07-11, gates green), receives the fused strip-2 kernel as experiment 1, compiles it, gates it, benchmarks 2,116.4 steps/s, and issues its own verdict:

```
{"exp": 1, "fused": true, "strip": 2,
 "compile_ok": true, "verify_ok": true,
 "steps_per_sec": 2116.4, "best_sps": 2116.4,
 "verdict": "keep"}
```

Then the driver tries to cheat: it submits the same kernel with rainfall injection doubled in one of the two cell updates, a change that compiles and breaks conservation. The gate fails it; the driver requests `benchmark` anyway, and the server answers:

```
{"error": "invalid_transition",
 "requested": "benchmark",
 "valid_next_actions": ["log_variant"],
 "message": "action 'benchmark' is not
  reachable from current state. Valid
  actions now: ['log_variant']."}

```

The variant is ledgered as `revert` with a null `steps_per_sec`.

How repeatable is this? `passk_flood2.py` runs the two-cycle reproduction `k` times from cool starts, scoring a pass when the baseline gates green inside 1,300 to 1,400 steps/s, the fused kernel is kept at 1.5x or better, and the broken kernel is refused and nulled. Five of five runs passed on 2026-08-09: baselines  $1,348.6 \pm 4.1$  steps/s, the kept kernel 2,127.7 in

TABLE II  
CONCURRENT GPU FLOOD AND CPU LLM DECODE

| Condition                    | GPU flood (steps/s)     | CPU decode (t/s) |
|------------------------------|-------------------------|------------------|
| Optimized kernel alone       | $2,128.0 \pm 1.7$ (n=3) |                  |
| Original kernel alone        | $1,351.4 \pm 0.5$ (n=3) |                  |
| Decode alone (soaked)        |                         | $11.4 \pm 0.0$   |
| Concurrent, optimized kernel | $883.4 \pm 47.1$ (n=28) | $9.6 \pm 0.2$    |
| Concurrent, original kernel  | $401.8 \pm 8.3$ (n=12)  | $9.6 \pm 0.2$    |

every run (identical at the timer’s resolution), speedups 1.574 to 1.585.

One session was fully agent-driven (2026-08-02, budget three). The server’s `characterize` payload names fixed per-invocation cost as the bound and exposes the `strip` parameter, so the session tests the loop rather than blind discovery: the model worked the exposed knob (strip 4 kept at 1,470.6, +8.8%), hit the register cliff at strip 8 (921.7, reverted), probed it at strip 6 (1,298.7, reverted), and never reached fusion; all of its reported numbers came from the machine.

## VII. CONCURRENT CPU AND GPU EXECUTION

In deployment the GPU loops the simulation while the CPU routes and decodes (Granite 4.0 1B, llama.cpp, `-ngl 0`, 4 threads; llama-bench prints its 1.63 B GGUF under a granite-3B size label). Decode, the token-by-token generation phase, is memory-bound, which matters because the two processors share one LPDDR bus.

The campaign turned up one configuration surprise first. With any Vulkan device visible, llama.cpp places CPU-resident model weights in the GPU’s host-pinned, write-combined memory even at `-ngl 0`, memory that is slow for the CPU reads that dominate decode. Hiding the device with `GGML_VK_VISIBLE_DEVICES=99` takes decode from 9.68 to 11.59 tokens/s, a 20% gain from an environment variable (A/B from cool starts, llama-bench `r=5`; raw logs in `docs/paper/artifacts/`, `ab_vkvisible`). Every number below uses the hidden-device configuration, which the application ships.

`pi/flood/conc_steady_bench.sh` measures each kernel alone (cool starts, finished before heat accumulates), decode alone, and each kernel looping while decode runs; concurrent and decode-alone phases are measured only after a decode soak reaches the firmware’s governed thermal plateau (about 76 C, soft-limit bit active), over windows of 12 to 28 flood completions (llama-bench `tg64`, `r=20`), counting a flood run only if it fits inside the window. An earlier short-window campaign, preserved in `conc_bench2/`, read 10 to 15% differently in both directions because it measured the thermal transient. Results from the steady-state run of 2026-08-10 are in Table II (raw logs in `docs/paper/artifacts/`, `conc_steady`).

Three things follow. Co-processing costs the CPU 16% of its decode and buys a continuous 883 steps/s of simulation that a CPU-only deployment does not have. The interference is lopsided: the CPU keeps 84% of its rate while the GPU

TABLE III  
THE CPU-ONLY COUNTERFACTUAL

| Condition                              | Flood (steps/s)           | CPU decode (t/s) |
|----------------------------------------|---------------------------|------------------|
| CPU flood alone, 1 / 2 / 4 threads     | 992.5 / 1,980.5 / 3,852.4 |                  |
| Decode alone, 3 threads (pinned)       |                           | 11.9 $\pm$ 0.0   |
| Partitioned: flood 1t + decode 3t      | 678.8 $\pm$ 10.6 (n=25)   | 8.4 $\pm$ 0.1    |
| Oversubscribed: flood 4t + decode 4t   | 857.1 $\pm$ 181.0 (n=89)  | 3.0 $\pm$ 0.3    |
| GPU concurrent, optimized (from above) | 883.4 $\pm$ 47.1          | 9.6 $\pm$ 0.2    |

keeps 42%, as decode traffic crowds the shared bus. And contention favors the optimized kernel: 1.58x over the original alone becomes 2.20x concurrent, at identical decode cost (9.6 t/s under either kernel).

Decode is also the GPU’s worst case. Paired with a compute-bound CPU load (openssl SHA-256, four threads) the GPU keeps 99% and the load keeps 88% of its 16 KB-block throughput; paired with the memory-bound four-thread CPU flood the GPU keeps 92% and the load keeps 75% (continuous sim-mode probes, both sides logged; docs/paper/artifacts/, genload2). The GPU is the robust side of every pairing except decode, which streams model weights from DRAM every token and drives GPU retention to 42%: interference tracks the CPU load’s memory traffic.

#### A. The CPU-Only Counterfactual

Four idle A76 cores run this stencil at 3,852 steps/s, well above the V3D’s 2,128, so the GPU looks unnecessary. pi/flood/cpuflood.cpp is the same update with OpenMP and passes the same three gates (NMSE 1.46e-11); pi/flood/cpu\_steady\_bench.sh measures it alone and sharing the cores with decode, pinned apart (flood on core 0, decode on 1 to 3) and oversubscribed (four threads each). The same soak-and-measure protocol applied; raw logs are in docs/paper/artifacts/, cpu\_steady. Table III presents the outcome.

The idle-core number is real but unavailable: in deployment, decode is always running. Oversubscribed, decode collapses 75% (llama.cpp’s workers synchronize every token; losing one core stalls all four), and the flood mean is inflated by bursts between decode repetitions, hence its  $\pm$  181. Partitioned, the best CPU-only arrangement, decode pays a 29% tax against the GPU split’s 16%; time-slicing (flooding about 23% of the time to match 883 steps/s) would cap decode at 8.8 t/s and freeze guidance during bursts. The GPU split beats the best measured CPU-only option on both axes, by 30% on simulation (883.4  $\pm$  47.1 vs 678.8  $\pm$  10.6) and 14% on decode (9.6  $\pm$  0.2 vs 8.4  $\pm$  0.1): the slower processor still wins, because the CPU has no spare cycles to sell.

## VIII. LIMITATIONS

The envelope is one board (Pi 5, V3D 7.1.10.2, Mesa v3dv 25.0.7) and one grid family, with speedups that shrink as the grid grows; headroom likely remains. The July and August campaigns straddle an OS update (kernel 6.12 to 6.18); flood baselines agree across it (1,346.8 before, 1,348.6  $\pm$  4.1 after). The gates are necessary rather than sufficient: verification covers one storm scenario at one grid size, and the keep decision rests on a single timing sample against a 1% threshold. The 3,803 steps/s CPU figure assumed a cache-resident 256x256 working set and may not survive larger grids. The scripted reproduction’s kernel came from the hand sweep, and the agent-driven session worked a knob the server itself exposed. The demonstrated property is machine verification; autonomous discovery remains open. Full GPU offload of the language model is blocked by an upstream llama.cpp Vulkan defect.

## IX. THE SAMPLE APPLICATION

The sample application [7] simulates surface flooding over real terrain [1], finds shelter routes by mobility profile [2], and explains the result in plain language, entirely offline; the model never makes a safety decision, only narrating what the deterministic code computed. The application lives at github.com/msradam/bonbibibi, distinct from this paper’s contributions; the harness at github.com/msradam/seppa. Measurements used its Granite 4.0 1B configuration (since moved to a larger model); build commands are in pi/flood/README.md, and cited references are archived in docs/paper/references/. Numbers without an archived log (the original harness’s 1,045 steps/s and the GEMM and matrix-vector results) come from the repository’s dated running notes.

## X. CONCLUSION

Seppa separates proposal from judgment: a language model suggests kernels for the Raspberry Pi 5’s integrated GPU, and a state machine the model cannot argue with decides what is true about them. On a real flood stencil this produced a verified 1.59x, reproduced within 1% by the machine five of five times, and it holds in deployment: 883 steps/s of simulation beside decode at 84% of its solo speed. What we trust in the end is not the model’s account of its work but the gate the work had to pass through. The suggestion is open-ended: commodity edge boards carry more usable silicon than their deployments exercise, and a loop that cannot lie about correctness is a reasonable way to get at it.

## REFERENCES

- [1] M. Guidolin, A. S. Chen, B. Ghimire, E. C. Keedwell, S. Djordjevic, and D. A. Savic. A weighted cellular automata 2D inundation model for rapid flood analysis. *Environmental Modelling & Software* 84:378-394, 2016.
- [2] J. Dibbelt, B. Strasser, and D. Wagner. Customizable Contraction Hierarchies. *ACM Journal of Experimental Algorithmics* 21, Article 1.5, 2016.
- [3] A. Ouyang, S. Guo, et al. KernelBench: Can LLMs Write Efficient GPU Kernels? arXiv:2502.10517, 2025.
- [4] X. Li et al. CUDA-L1: Improving CUDA Optimization via Contrastive Reinforcement Learning. arXiv:2507.14111; ICLR 2026.
- [5] Burr. Apache Software Foundation (incubating). <https://github.com/apache/burr>

- [6] Model Context Protocol. <https://modelcontextprotocol.io>
- [7] Bonbibì: offline flood-aware routing on a Raspberry Pi 5. <https://github.com/msradam/bonbibì>
- [8] llama.cpp. ggml-org. <https://github.com/ggml-org/llama.cpp>
- [9] V3D driver documentation, Mesa 3D. <https://docs.mesa3d.org/drivers/v3d.html>
- [10] J. Blocklove, S. Garg, R. Karri, and H. Pearce. Chip-Chat: Challenges and Opportunities in Conversational Hardware Design. ACM/IEEE Workshop on Machine Learning for CAD (MLCAD), 2023.
- [11] Z. Zhou et al. Edge Intelligence: Paving the Last Mile of Artificial Intelligence With Edge Computing. Proc. IEEE 107(8), 2019.
- [12] S. Boddu and A. Mukherjee. Efficient Edge Deployment of Quantized YOLOv4-Tiny for Aerial Emergency Object Detection on Raspberry Pi 5. arXiv:2506.09300, 2025.
- [13] A. Reuther et al. Survey of Machine Learning Accelerators. IEEE HPEC, 2020.
- [14] The White House. Building Resilient Supply Chains, Revitalizing American Manufacturing, and Fostering Broad-Based Growth: 100-Day Reviews under Executive Order 14017. June 2021.
- [15] AutoKernel. RightNow AI. <https://github.com/RightNow-AI/autokernel>